



CPC357 ASSIGNMENT 2

IOT ARCHITECTURE AND SMART APPLICATIONS

SCHOOL OF COMPUTER SCIENCE

UNIVERSITI SAINS MALAYSIA

SEMESTER 1 2024/2025

ASSIGNMENT 2

Lecturer: Dr. Mohammad Nadhir Ab Wahab

Group Members:

Name	Matric No
BRYON SAVERO A/L MICHAEL LEO	159351
SUARNAA MUGUNTHAN A/L RAMAN	159208

Submission: 26th January 2025

Table Of Contents

1. Introduction	3
2. High-Level System Architecture Overview	5
2.1 Hardware Layer.....	6
2.2 Connectivity and Communication Layer	6
2.3 Cloud Infrastructure Layer	7
2.4 Data Management Layer	7
2.5 Visualization Layer	8
3. Design Considerations.....	10
3.1 Factors Influencing Design Decisions.....	10
1. Sensor Selection	10
2. Hardware Selection.....	10
3. Scalability	10
4. Cost-Effectiveness	10
5. Data Security	11
6. Connectivity and Communication	11
3.2 Trade-Offs	11
Cost vs. Performance	11
Ease of Use vs. Security	12
Real-Time Data vs. Long-Term Analysis	12
Flexibility vs. Specialization	12
4. Development Process	13
4.1 Setting Up Google Cloud Platform Services.....	13
Compute Engine	13
Firewall Rule	16
Service Account.....	17
BigQuery	19
4.2 Setting Up Hardware	20
Arduino Code	20
4.3 Setting Up Environment in Google VM	21
4.4 Creating the python script and running it.....	21
5. Security Considerations.....	23
5.1 Security Measures Implemented	23
5.2 Strategies for Data Protection and Access Control	24
6. Appendix.....	27

1. Introduction

The **weather monitoring station** is an **Internet of Things (IoT)**-based application designed to measure and report **environmental conditions** in **real time**. This system is built to provide **accurate** and **reliable data** on key environmental parameters, specifically **temperature**, **humidity**, and **rainfall**. Such data is crucial for a wide range of applications, from **agricultural planning** to **urban development**. Leveraging the widely used **DHT11 sensor** for temperature and humidity measurements and a **rain sensor** for detecting precipitation, the station offers a **cost-effective** and **efficient solution** for continuous environmental monitoring.

By integrating these sensors with **Google Cloud Platform (GCP)**, the system achieves several critical objectives: **scalability**, **real-time data processing**, and **secure data storage**. The integration allows the collected data to be transmitted seamlessly from **physical sensors** to the **cloud**, enabling efficient **storage**, **analysis**, and **visualization**. The use of **GCP** ensures the system can handle **increased data volumes** and provide reliable access to the **information** it generates.

This project addresses the growing need for **efficient** and **reliable environmental monitoring solutions**. As the world faces challenges such as **climate change** and **urban expansion**, monitoring systems like this play a critical role in supporting **data-driven decision-making**. For instance, in **agriculture**, accurate weather data helps farmers optimize **planting schedules**, **irrigation**, and **harvests**. In **urban planning**, such systems can support **drainage design** and **flood prevention** measures. Moreover, in **disaster management**, **real-time weather data** can improve **preparedness** and **response strategies**, potentially saving **lives** and **resources**.

The project emphasizes the interplay between **hardware** and **software components**, ensuring **seamless communication** and **robust data handling**. The **hardware layer** collects environmental data from **sensors**, while the **software layer** processes, stores, and analyzes this data using **cloud-based tools**. This **dual-layered approach** ensures that the system operates **efficiently** and **reliably**. The design also focuses on **flexibility** and **scalability**, enabling the addition of **new sensors** or functionalities without disrupting the existing setup. This **forward-looking architecture** ensures the weather monitoring station can adapt to **future technological advancements** and **emerging use cases**.

By focusing on these principles, the **weather monitoring station** exemplifies how **IoT technologies** can bridge the gap between **physical** and **digital systems**. It provides a foundation for building **smarter**, more **sustainable environments**. The system's ability to deliver **real-time, actionable insights** makes it a valuable tool across various domains, reinforcing the importance of **environmental monitoring** in today's **data-driven world**.

2. High-Level System Architecture Overview

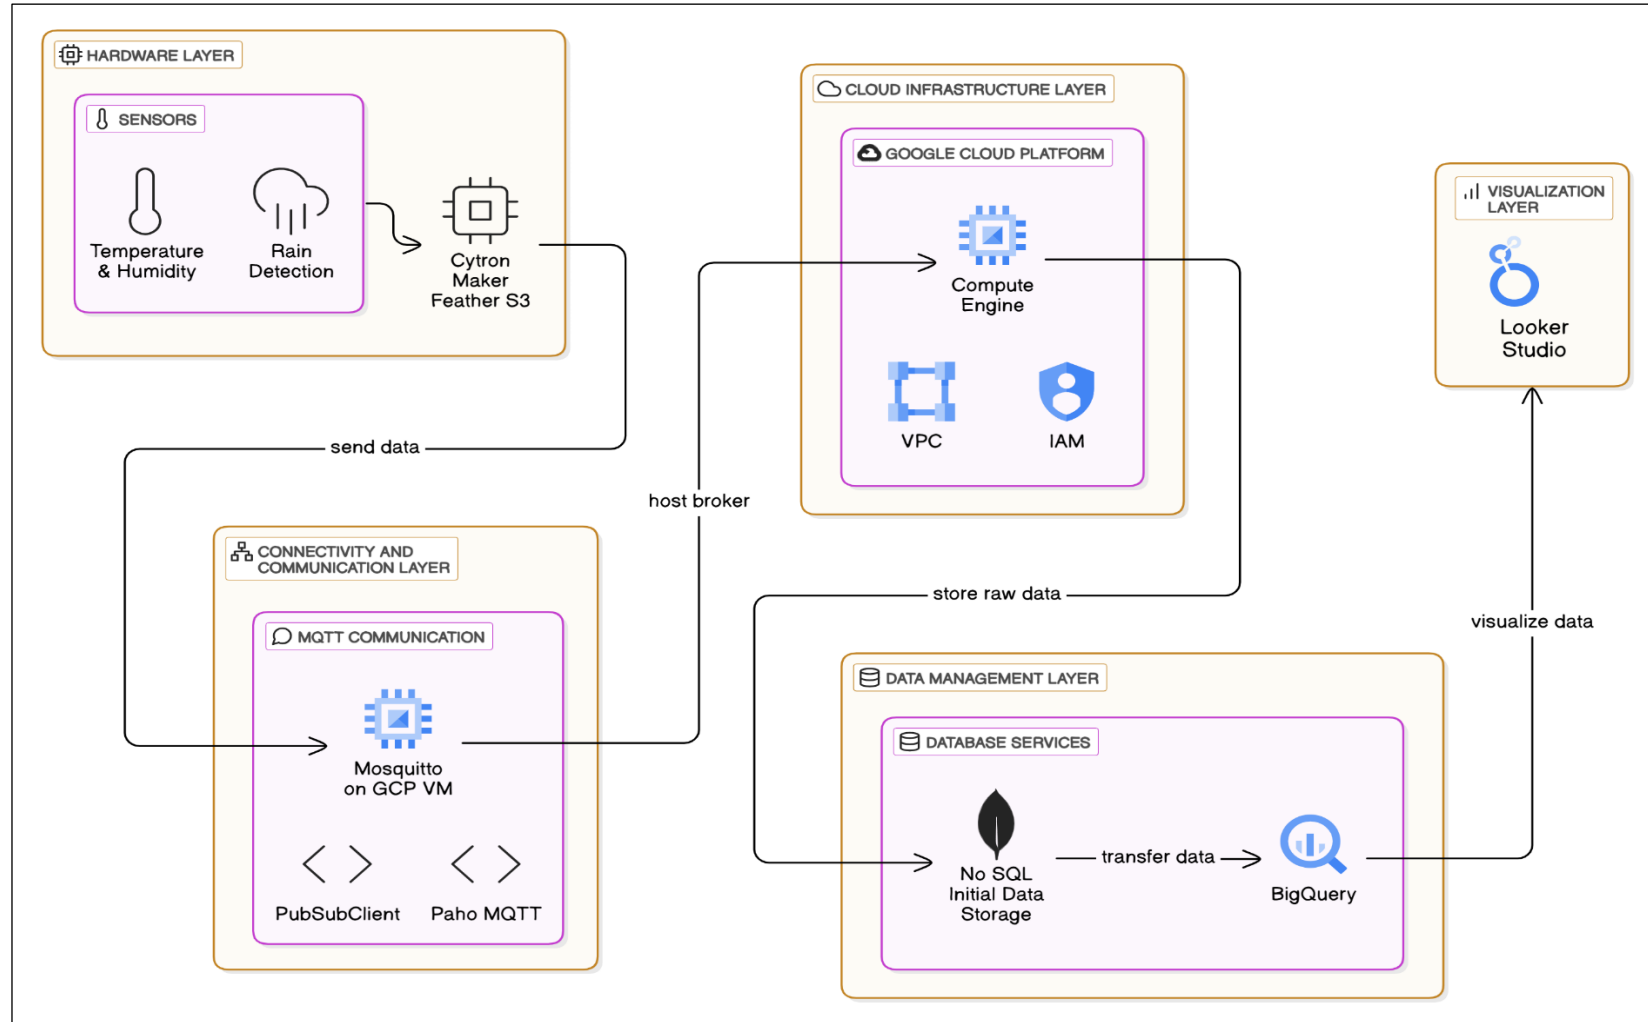


Figure 1: High-level architecture diagram of the system

2.1 Hardware Layer

The hardware layer forms the foundational component of the weather monitoring system and comprises the physical IoT devices responsible for environmental data collection. At the core of this layer is the **Cytron Maker Feather S3 microcontroller**, which functions as the primary data acquisition and transmission unit. This microcontroller is equipped with two sensors to measure environmental parameters:

- **DHT11 Sensor:** This sensor is responsible for measuring temperature and humidity levels in the environment. It provides digital output and is known for its simplicity and cost-effectiveness, making it suitable for IoT applications.
- **Rain Sensor:** This sensor detects precipitation levels and provides real-time data regarding rainfall. It can determine the presence and intensity of rain, enabling the system to monitor weather conditions effectively.

The microcontroller collects the raw environmental data from these sensors and transmits it to the cloud for further processing and analysis. The selection of reliable and efficient hardware components ensures accurate data collection, forming the basis for the system's functionality.

2.2 Connectivity and Communication Layer

The connectivity and communication layer is a critical component responsible for the secure and efficient transmission of sensor data from the IoT devices to the cloud infrastructure. This layer leverages the following key technologies:

- **Mosquitto MQTT Broker:** The Mosquitto MQTT broker is hosted on a Google Cloud Platform (GCP) Compute Engine virtual machine. It serves as the central messaging hub, managing the publication and subscription of data between IoT devices and the server.
- **MQTT Protocol:** The **MQTT (Message Queuing Telemetry Transport)** protocol is used for lightweight and efficient data transmission. This protocol is particularly well-suited for IoT applications due to its low bandwidth requirements. For secure communication, **TLS (Transport Layer Security)** is employed on port **8883**, ensuring data integrity and confidentiality during transmission.
- **PubSubClient Library:** On the microcontroller side, the **PubSubClient library** is utilized to enable MQTT communication. It allows the microcontroller to publish the collected sensor data to the broker.

- **Paho MQTT Library:** On the server side, the **Paho MQTT library** facilitates interaction with the MQTT broker, enabling the reception of incoming sensor data for further processing.

This robust communication framework ensures seamless and secure data transfer between the IoT devices and the cloud infrastructure, forming the backbone of the system's operation.

2.3 Cloud Infrastructure Layer

The cloud infrastructure layer is designed to host and manage the essential services required for the weather monitoring solution. It leverages several Google Cloud Platform (GCP) services to ensure scalability, security, and reliability:

- **Compute Engine:** The Mosquitto MQTT broker runs on a **Compute Engine virtual machine**, which provides a high-performance and customizable computing environment.
- **Virtual Private Cloud (VPC):** The **VPC service** is used to configure a secure and isolated network for the system. It manages the network communication and enforces **firewall rules** to protect the MQTT broker and other cloud components from unauthorized access.
- **Identity and Access Management (IAM):** The **IAM service** ensures appropriate access control and permissions for the various system components. A service account has been specifically created for **BigQuery** with the role of **BigQuery Admin**, ensuring proper permissions for managing and querying the data warehouse. This enhances security by restricting access to only authorized users and applications.

The infrastructure layer provides a scalable and secure environment to support the system's requirements by leveraging these cloud services, ensuring uninterrupted data processing and storage.

2.4 Data Management Layer

The data management layer focuses on the storage and processing of sensor data collected from the IoT devices. This layer incorporates the following components:

- **MongoDB:** The raw sensor data is initially stored in a **MongoDB database** hosted on the Compute Engine virtual machine. MongoDB, being a NoSQL database, offers flexibility in handling unstructured or semi-structured data. It serves as a temporary storage solution, enabling efficient ingestion of real-time data.
- **BigQuery:** The collected data is subsequently transferred from MongoDB to **BigQuery**, a fully managed, highly scalable data warehouse service provided by GCP. BigQuery enables fast and efficient analysis of large datasets, making it ideal for processing and analysing the high volume of sensor data generated by the system.

This layered data management approach ensures that the system can handle both the real-time ingestion of sensor data and the long-term storage and analysis of large datasets.

2.5 Visualization Layer

The visualization layer is responsible for presenting the processed sensor data in an intuitive and insightful manner. This is achieved using Looker Studio, a powerful business intelligence tool offered by GCP:

- **Interactive Dashboards:** Looker Studio is utilized to create interactive dashboards that allow users to visualize and analyze the weather-related data effectively. These dashboards provide a user-friendly interface to explore trends, patterns, and insights derived from the sensor data.
- **Seamless Integration with BigQuery:** Looker Studio integrates seamlessly with BigQuery, enabling real-time access to the processed data stored in the data warehouse. This integration facilitates the creation of dynamic visualizations that update automatically as new data becomes available.

The system provides stakeholders with actionable insights and a comprehensive understanding of weather conditions, enhancing decision-making capabilities by utilizing Looker Studio.

This weather monitoring station architecture demonstrates a comprehensive and well-structured IoT solution that leverages various Google Cloud Platform (GCP) services to collect, store, process, and visualize environmental data. By adopting a layered approach—consisting of distinct hardware, connectivity, cloud infrastructure, data management, and visualization components—the system ensures scalability, security, and flexibility. This architecture can

provide valuable weather-related insights, making it a robust and efficient solution for monitoring environmental conditions.

3. Design Considerations

3.1 Factors Influencing Design Decisions

1. Sensor Selection

- The **DHT11 sensor** was chosen for its **affordability**, **simplicity**, and **adequate precision** for measuring temperature and humidity. It provides **digital output**, making it ideal for IoT applications where low-cost and reliable data collection is prioritized. While the DHT11 may lack the accuracy of advanced sensors like the DHT22, it offers sufficient performance for general weather monitoring tasks, ensuring a cost-effective solution.
- The **rain sensor** was selected for its ability to provide both **digital** and **analog outputs**, enabling it to detect the **presence** and **intensity** of rainfall. While it may not offer the precision of advanced rain gauges, its versatility and cost-effectiveness make it suitable for the system's real-time monitoring requirements.

2. Hardware Selection

- The **Cytron Maker Feather S3 microcontroller** was chosen due to its **Wi-Fi** and **Bluetooth capabilities**, which facilitate seamless connectivity with cloud services. Its **compatibility with MQTT** protocols and support for lightweight data processing ensures efficient operation. Compared to less advanced microcontrollers, this choice balances performance and energy efficiency, making it ideal for IoT applications like this weather monitoring station.

3. Scalability

- The system's architecture supports **future expansion**, allowing for the integration of additional sensors, such as **wind speed detectors** or **air quality monitors**, without requiring significant modifications. The use of **Google Cloud Platform (GCP)** services, such as **BigQuery** and **Looker Studio**, ensures the system can handle **large data volumes** and **dynamic visualizations**. This scalability makes the design flexible for future requirements or changing application needs.

4. Cost-Effectiveness

- The integration of **low-cost sensors** like the DHT11 and rain sensor, combined with GCP's **pay-as-you-go model**, ensures the system remains affordable for a wide range of users. While components like MongoDB offer efficient real-

time data ingestion, and BigQuery handles large-scale analytics, the system avoids unnecessary expenditure by leveraging managed cloud services and cost-effective hardware. While more advanced sensors and cloud features might improve performance, the chosen components and configurations provide a **balanced approach** to maintaining affordability.

5. Data Security

- The system implements several **security measures** to protect data integrity and confidentiality. **Transport Layer Security (TLS)** is used for MQTT communication, ensuring secure data transmission. The **Virtual Private Cloud (VPC)** enforces **firewall rules**, while **Identity and Access Management (IAM)** provides **role-based access control (RBAC)**, restricting access to only authorized users and applications. These measures ensure compliance with IoT security best practices, even though they introduce additional setup complexity.

6. Connectivity and Communication

- The use of **Mosquitto MQTT Broker** on a **GCP Compute Engine VM** provides a reliable and centralized messaging hub. The lightweight **MQTT protocol**, combined with libraries like **PubSubClient** (microcontroller-side) and **Paho MQTT** (server-side), ensures **efficient communication** between hardware and cloud infrastructure. This robust framework minimizes latency and bandwidth usage, which are critical for IoT systems.

3.2 Trade-Offs

Cost vs. Performance

- A deliberate decision was made to use **cost-effective sensors** like the DHT11 and rain sensor, even though they may lack the precision of more expensive alternatives. For instance, while the DHT11 is less accurate than the DHT22, it provides sufficient data accuracy for **general monitoring purposes**. Similarly, the rain sensor offers basic precipitation detection but cannot provide **quantitative rainfall measurements** like advanced rain gauges. These trade-offs were necessary to achieve the project's objective of creating a **budget-friendly solution** without compromising the system's primary functionality.

Ease of Use vs. Security

- Implementing **advanced security configurations**, such as **strict firewall rules** and **encrypted MQTT communication**, added complexity to the system's deployment and maintenance. These measures were essential for ensuring **robust data protection**, but they required additional setup and expertise, potentially increasing the learning curve for users. The decision to prioritize security over ease of use reflects the importance of **data integrity** and **user trust** in IoT applications. This trade-off ensures that the system remains **secure and reliable**, even if it slightly reduces accessibility for non-technical users.

Real-Time Data vs. Long-Term Analysis

- The architecture prioritizes real-time data ingestion using **MongoDB** while leveraging **BigQuery** for long-term storage and analytics. While this layered approach ensures efficiency and scalability, it introduces an additional step in the data pipeline, slightly increasing system complexity. This trade-off is justified by the system's ability to handle both **real-time operations** and **historical data analysis** effectively.

Flexibility vs. Specialization

- The decision to use general-purpose components like the **Maker Feather S3** and **rain sensor** emphasizes flexibility but sacrifices some precision and specialization. For example, using a **dedicated weather station kit** could simplify integration but reduce the ability to customize the system for unique use cases.

4. Development Process

4.1 Setting Up Google Cloud Platform Services

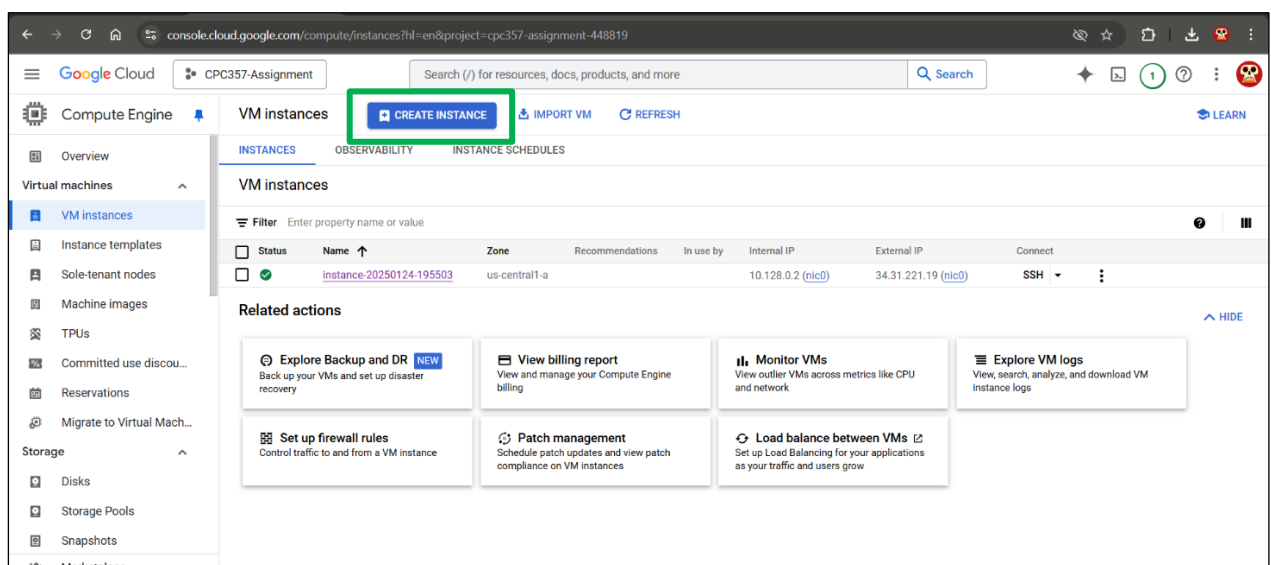
Services that are used for the development:

- Compute Engine (VM Instances)
- VPC Network (Firewall Rule)
- IAM & Admin (Service Accounts)
- BigQuery
- Looker Studio

Compute Engine

Compute Engine is a customizable compute service that allows you to create and run virtual machines on Google Cloud. We can opt in for pre-built and ready-to-go configurations of VMs for a quick start, or can go for custom machine types and create VMs with the optimal amount of processing power and memory based on our needs. Go to this link to navigate to the Google Cloud Console. If you are using Compute Engine for the first time, you need to complete a few additional steps before you can proceed with the rest of the session. Please create a project in the console and enable the Compute Engine API. Note that you also need to set up a billing account in order to use the service.

Next, go to the navigation menu and select **Compute Engine > Virtual Machines > VM Instances**. Click the **Create Instance** button to create a new virtual machine.



Under **Machine configuration**, you can set the region and zone that runs your virtual machine. Next, go to **OS and storage** in the navigation menu to choose the operating system. Click the Change button and select **Ubuntu** for the operating system and **Ubuntu 20.04 LTS** for the version.

Machine configuration

Name *
Instance-20250126-133230

Region *
us-central1 (Iowa)
Region is permanent

Zone *
Any
Google will choose a zone on your behalf, maximizing VM obtainability. Zone is permanent.

☒ General purpose ☐ Compute optimized ☐ Memory optimized ☐ Storage optimized ☐ GPUs

Machine types for common workloads, optimized for cost and flexibility

Series	Description	vCPUs	Memory	Platform
<input type="radio"/> C4	Consistently high performance	2 - 192	4 - 1,488 GB	Intel Emerald Rapids
<input type="radio"/> C4A	Arm-based consistently high performance	1 - 72	2 - 576 GB	Google Axion
<input type="radio"/> N4	Flexible & cost-optimized	2 - 80	4 - 640 GB	Intel Emerald Rapids
<input type="radio"/> C3	Consistently high performance	4 - 192	8 - 1,536 GB	Intel Sapphire Rapids
<input type="radio"/> C3D	Consistently high performance	4 - 360	8 - 2,880 GB	AMD Genoa
<input checked="" type="radio"/> E2	Low cost, day-to-day computing	0.25 - 32	1 - 128 GB	Based on availability
<input type="radio"/> N2	Balanced price & performance	2 - 128	2 - 864 GB	Intel Cascade Lake
<input type="radio"/> N2D	Balanced price & performance	2 - 224	2 - 896 GB	AMD EPYC
<input type="radio"/> T2A	Scale-out workloads	1 - 48	4 - 192 GB	Ampere Altra ARM

Monthly estimate
\$25.46
That's about \$0.03 hourly

Pay for what you use: no upfront costs and per second billing

Item	Monthly estimate
2 vCPU + 4 GB memory	\$24.46
10 GB balanced persistent disk	\$1.00
Total	\$25.46

[Compute Engine pricing](#)

[LESS](#)

CREATE **CANCEL** [EQUIVALENT CODE](#)

Operating system and storage

Name
Instance

Type
New boot disk

Size
10 GB

Snapshot schedule
No schedule

License type
Free

Image
Debian GNU/Linux 12 (bookworm)

CHANGE

Additional storage and VM back

Backup plan **NEW**

Secure your backups against deletion through Backup and DR Service, a separate service. [Learn more](#)

Backup plan

Container

Boot disk

Select an image or snapshot to create a boot disk, or attach an existing disk. Can't find what you're looking for? Explore hundreds of VM solutions in [Marketplace](#)

PUBLIC IMAGES **CUSTOM IMAGES** **SNAPSHOTS** **ARCHIVE SNAPSHOTS** **EXISTING DISKS**

Operating system
Ubuntu

Version *
Ubuntu 20.04 LTS

x86_64, amd64 focal image built on 2025-01-11

Boot disk type *
Balanced persistent disk

COMPARE DISK TYPES

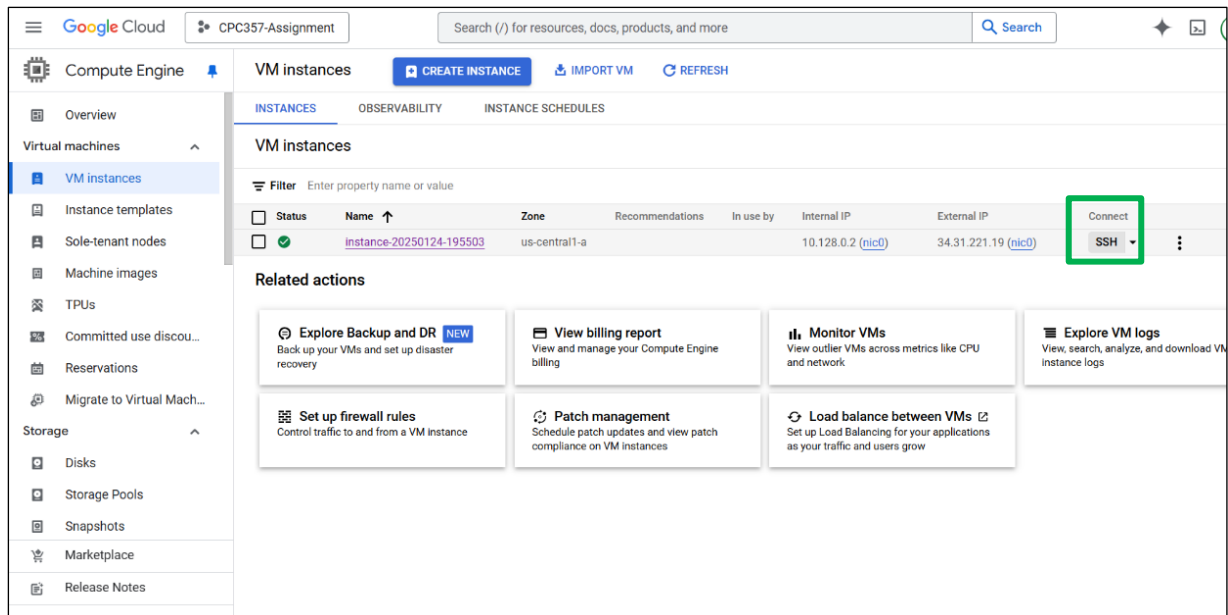
Size (GB) *
10

Provision between 10 and 65536 GB

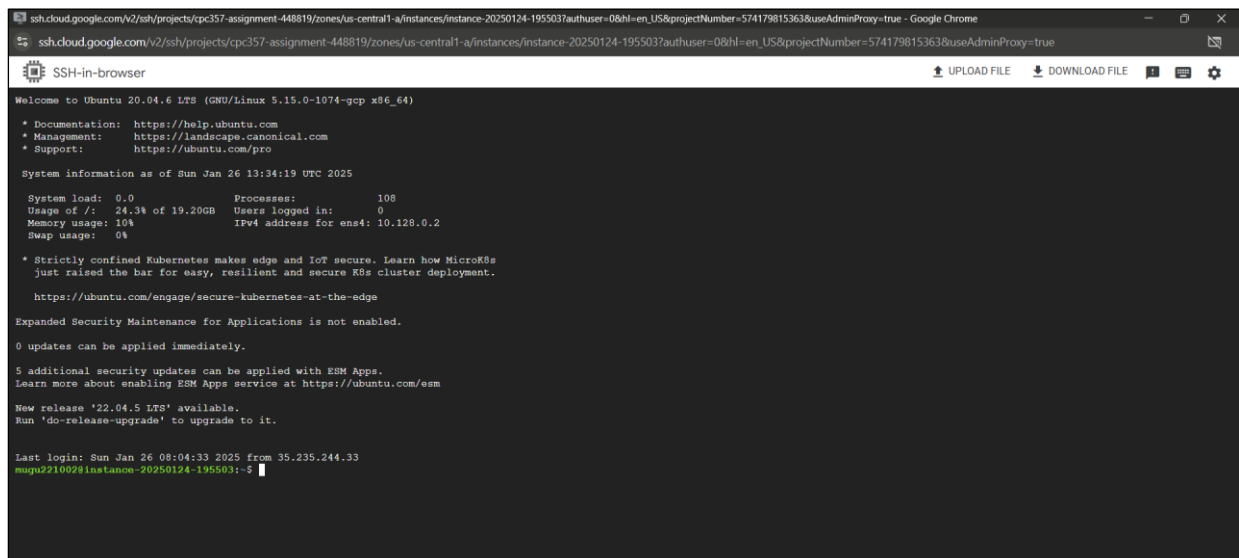
SHOW ADVANCED CONFIGURATION

SELECT **CANCEL**

Click the **Create** button to finish up the instance creation. Upon completion, you should be able to find your newly created VM instance in the list. Next, click the **SSH** button next to the VM instance to open up the SSH-in-browser.



A pop-up window will appear and click the **Authorize** button to give permission for the SSH-in-browser to connect to the VM.



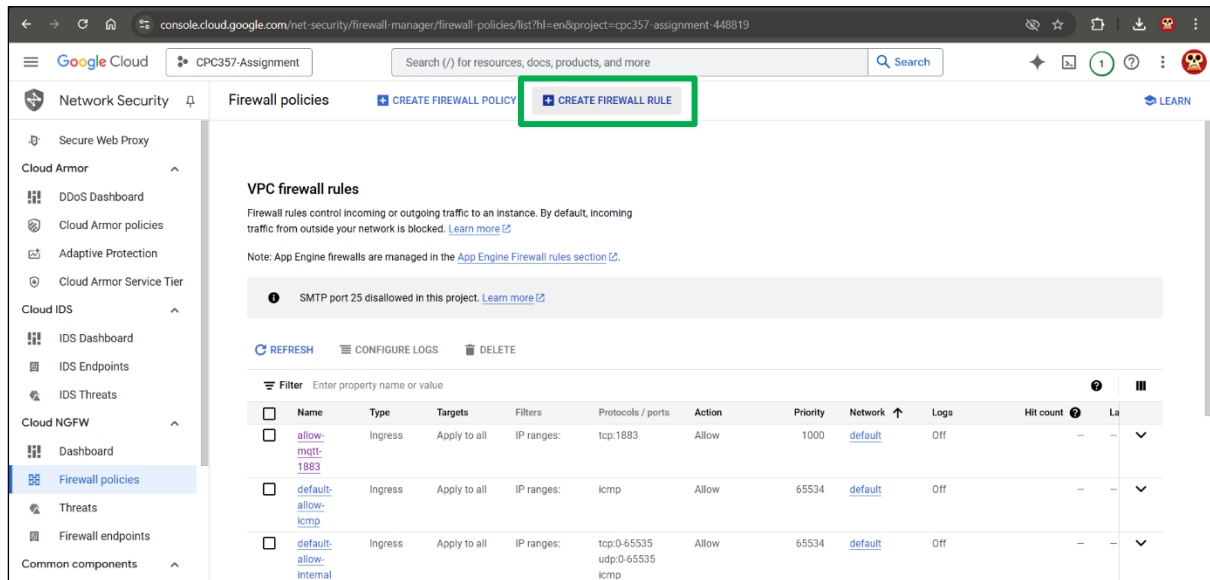
After that, run the following commands in the SSH-in-browser to setup the VM:

```
$ sudo apt-get update
```

```
$ sudo apt-get upgrade
```

Firewall Rule

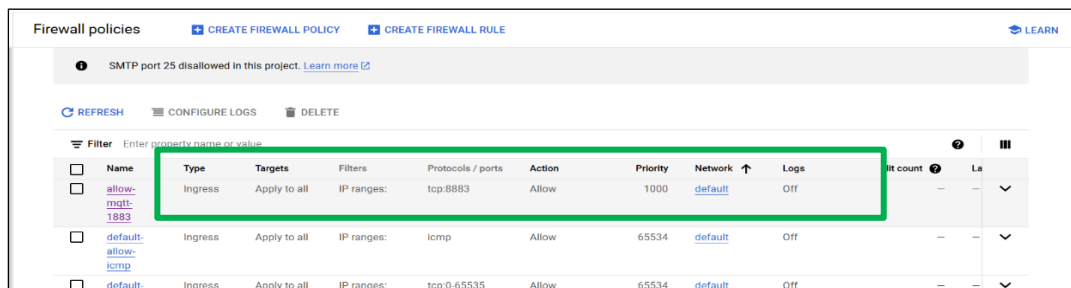
To provide connectivity for your Compute Engine VM instances, you need to configure the Virtual Private Cloud (VPC) network and set up the VPC firewall rules accordingly. To put it simply, VPC firewall rules let you allow or deny traffic to or from VM instances in a VPC network based on port number, tag, or protocol. First, go to the navigation menu and select **VPC Network > Firewall**. Then, click the **Create Firewall Rule** button at the top.



Configure the VPC firewall rule according to the specifications below:

- Name: allow-mqtt-8883
- Description: Allow traffic on port 8883
- Logs: off
- Network: default
- Priority: 1000
- Direction of traffic: Ingress
- Action on match: Allow
- Targets: All instances in the network
- Source filter: IPv4 ranges
- Source IPv4 ranges: 0.0.0.0/0
- Second source filter: None
- Destination filter: None
- Protocols and ports: Specified protocols and ports (TCP Port 8883)

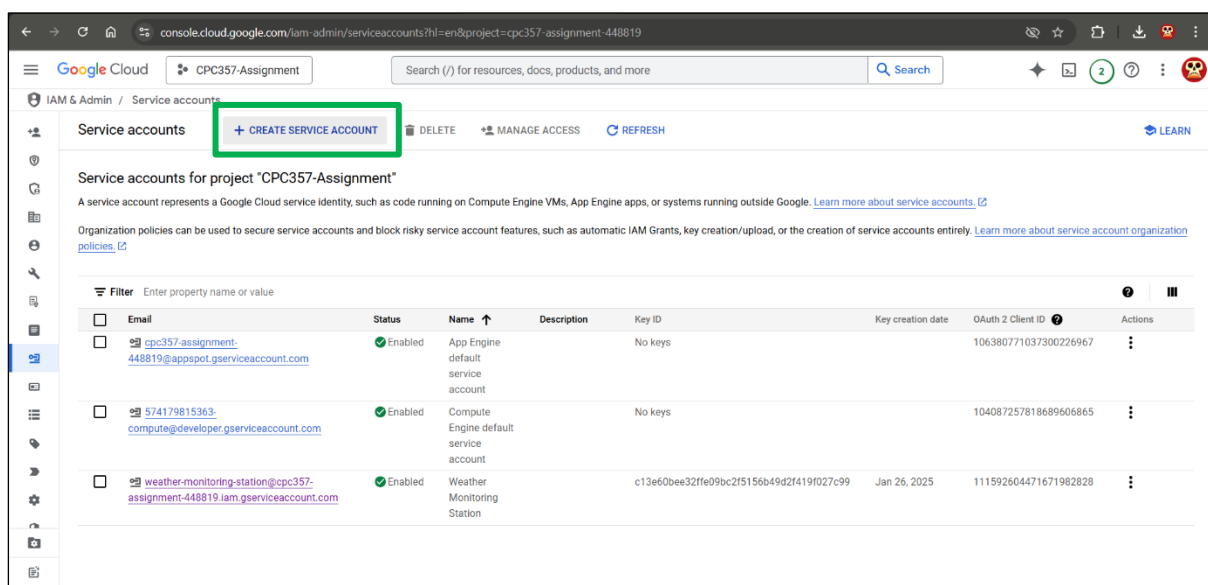
Once you're done, click on the Create button to save and create the VPC firewall rule and upon completion, the newly created VPC firewall rule will be displayed in the list.



Name	Type	Targets	Filters	Protocols / ports	Action	Priority	Network	Logs	Hit count	Last modified
allow-mysql-1883	Ingress	Apply to all	IP ranges:	tcp:8883	Allow	1000	default	Off	—	—
default-allow-icmp	Ingress	Apply to all	IP ranges:	icmp	Allow	65534	default	Off	—	—
default	Ingress	Apply to all	IP ranges:	tcp:0-65535	Allow	65534	default	Off	—	—

Service Account

A service account represents a Google Cloud service identity, such as code running on Compute Engine VMs, App Engine apps, or systems running outside Google. Organization policies can be used to secure service accounts and block risky service account features, such as automatic IAM Grants, key creation/upload, or the creation of service accounts entirely. Go to the navigation menu and select **IAM & Admin > Service Accounts**. Then, click the **Create Service Account** button at the top.



Email	Status	Name	Description	Key ID	Key creation date	OAuth 2.0 Client ID	Actions
cpc357-assignment-448819@appspot.gserviceaccount.com	Enabled	App Engine default service account		No keys		106380771037300226967	⋮
574179815363-compute@developer.gserviceaccount.com	Enabled	Compute Engine default service account		No keys		104087257818689606865	⋮
weather-monitoring-station@cpc357-assignment-448819.iam.gserviceaccount.com	Enabled	Weather Monitoring Station		c13e60bee32ffe09bc2f5156b49d2f419f027c99	Jan 26, 2025	111592604471671982828	⋮

Under **Service account details**, give your account a name and click **Create and Continue**. Then, grant a role to your service account which is **BigQuery Data Editor** (access to edit all the contents of datasets). After that, create the service account by clicking **Done**.

Google Cloud | CPC357-Assignment | Search (/) for resources, docs, g

IAM & Admin / Service accounts / Create service account

Create service account

- Service account details**

Service account name:

Display name for this service account:

Service account ID:

Email address: weather-monitoring-station-404@cpc357-assignment-448819.iam.gserviceaccount.com

Service account description:

[CREATE AND CONTINUE](#)
- Grant this service account access to project (optional)**
- Grant users access to this service account (optional)**

[DONE](#) [CANCEL](#)

Google Cloud | CPC357-Assignment | Search (/) for resources, docs, g

IAM & Admin / Service accounts / Create service account

Create service account

- Service account details**
- Grant this service account access to project (optional)**

Grant this service account access to CPC357-Assignment so that it has permission to complete specific actions on the resources in your project. [Learn more](#)

Role: IAM condition (optional) [+ ADD IAM CONDITION](#)

Access to edit all the contents of datasets

[+ ADD ANOTHER ROLE](#)

[CONTINUE](#)
- Grant users access to this service account (optional)**

[DONE](#) [CANCEL](#)

After creating the service account, the account will be displayed in the list. Click the account and navigate to **Keys** and click **Add Key > Create New Key**. Then, in the pop-up window, choose **JSON** as the key type and click **Create**. After clicking, a JSON file containing the service account's key will be downloaded automatically. Keep the file safe because we will use that key in our python script later.

Google Cloud | CPC357-Assignment | Search (/) for resources, docs, products, and more

IAM & Admin / Service accounts / Service account: 100716468920290961819 / Keys

weather monitoring station

DETAILS | PERMISSIONS | **KEYS** | METRICS | LOGS

Keys

Service account keys could pose a security risk if compromised. [Learn more about service accounts on Google Cloud](#)

Google automatically disables service account keys detected as compromised.

Add a new key pair or upload a public key certificate from an existing key.

Block service account key creation using organization policies. [Learn more about setting organization policies for service accounts](#)

[ADD KEY](#)

Type	Status	Key	Creation date	Expiration date
No rows to display				

Create private key for "weather monitoring station"

Downloads a file that contains the private key. Store the file securely because this key can't be recovered if lost.

Key type

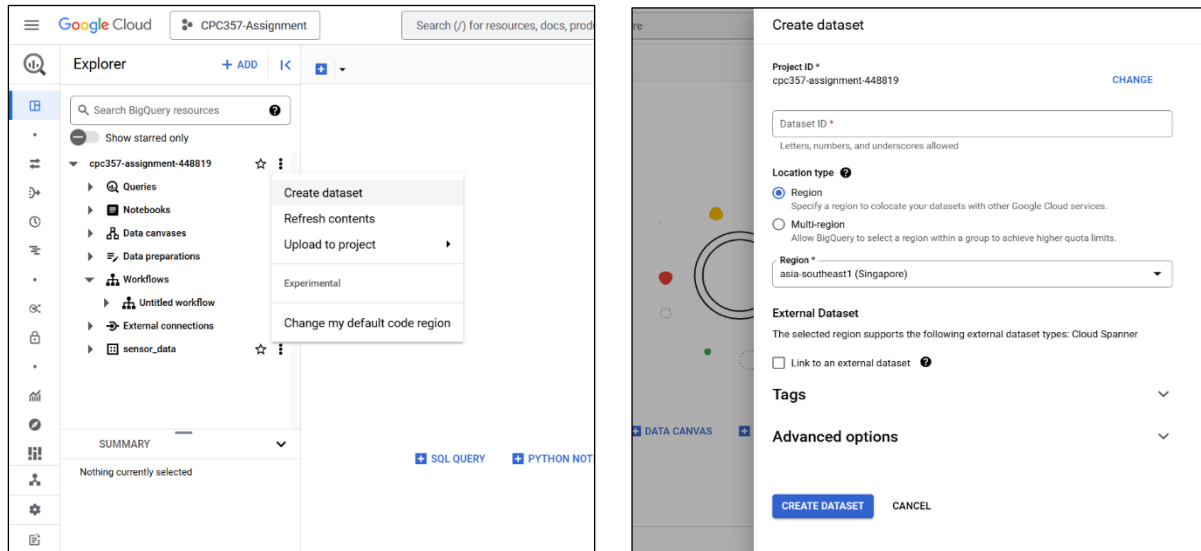
☒ JSON
Recommended

☐ P12
For backward compatibility with code using the P12 format

[CANCEL](#) [CREATE](#)

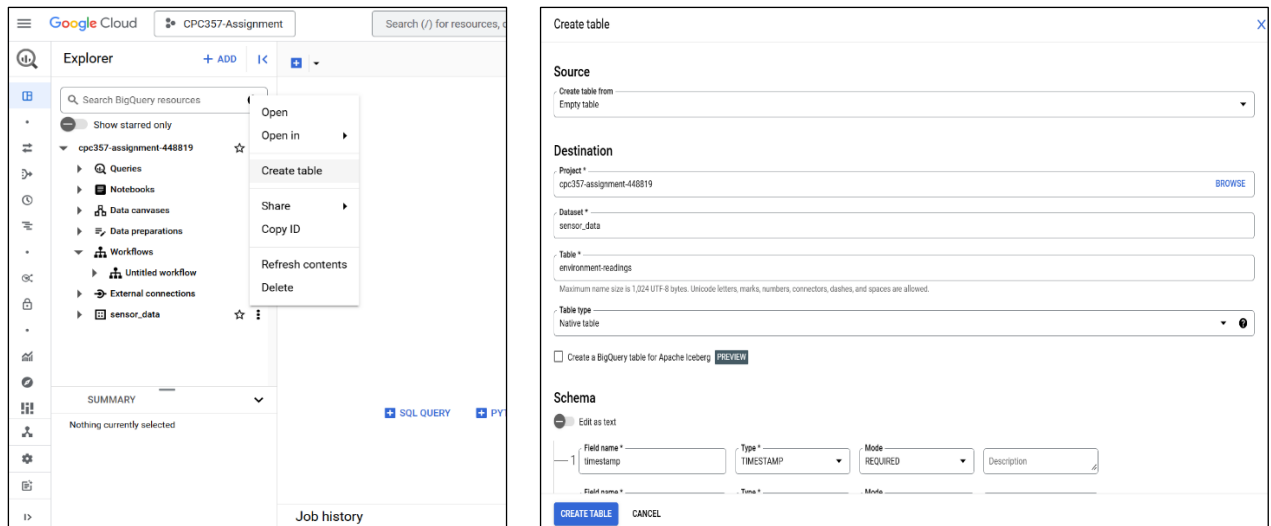
BigQuery

BigQuery is a serverless, cost-effective, and multicloud data warehouse designed to help you turn big data into valuable business insights. Navigate to BigQuery and click **Create dataset** under your project-id. Then, give enter your **Dataset ID** and specify the location type that you want to use to store your data. After that, click **Create**.

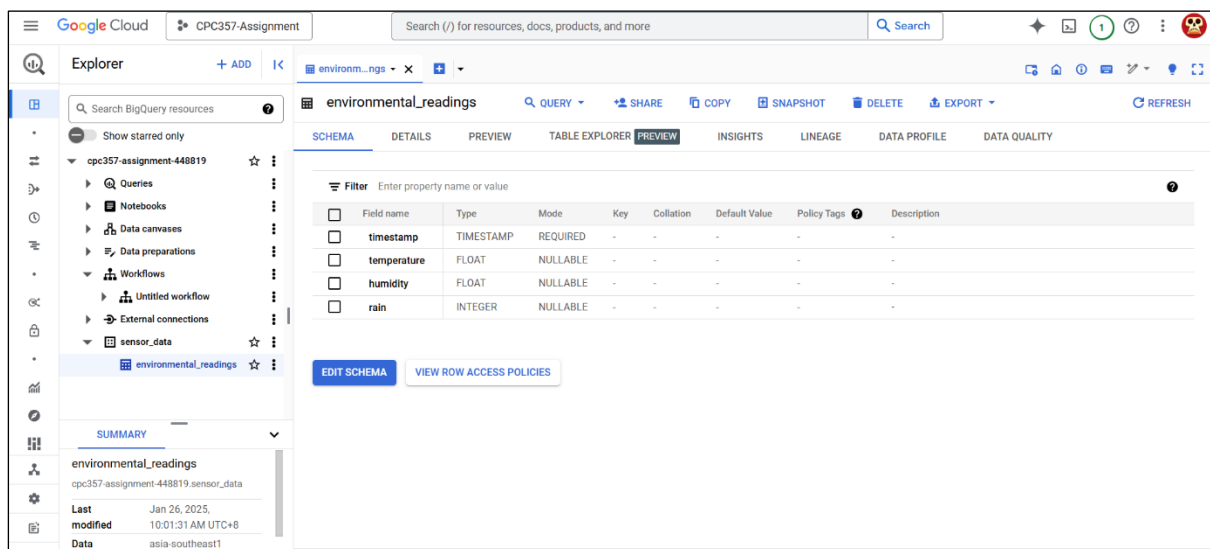


Upon completion, **Create table** under the newly created dataset and give table a name. After that, create four fields as below table under the **Schema** section:

Field name	Type	Mode
timestamp	TIMESTAMP	REQUIRED
temperature	FLOAT	NULLABLE
humidity	FLOAT	NULLABLE
rain	INTEGER	NULLABLE



Then click **Create Table**. Now you can see the created table with the column names and their types.



4.2 Setting Up Hardware

Arduino Code

Copy the code from the GitHub Repo: <https://github.com/mugunthantypical/CPC357-Assignment-2/blob/main/arduino-code/arduino-code.ino> and paste it in your Arduino IDE. Before uploading into your microcontroller, make sure you download and enable **PubSubClient library** in Arduino IDE.

4.3 Setting Up Environment in Google VM

Run all the following commands before creating the python script:

```
$ sudo apt-get install mosquitto
```

```
$ sudo apt-get install mosquitto-clients
```

```
$ sudo apt install python3-pip
```

```
$ pip install paho-mqtt
```

```
$ sudo apt-get install -y mongodb
```

```
$ pip install pymongo
```

```
$ pip3 install pymongo google-cloud-bigquery paho-mqtt
```

4.4 Creating the python script and running it

Use the python script from GitHub repo: <https://github.com/mugunthantypical/CPC357-Assignment-2/blob/main/mongo-bq.py>.

Simple follow the below command to create and run the script.

```
$ nano mongo-bq.py
```

Then copy the code and paste into the nano editor. After that, click **Ctrl+O** to save and **Ctrl+X** to exit the nano editor. You have to include your service account key in the same directory too. Copy the content from the JSON file , create a **service-account-key.json** file using the same command that used to create the python script, and paste it into the file. Save and close the file then run the below command to execute python script.

```
$ python3 mongo-bq.py
```

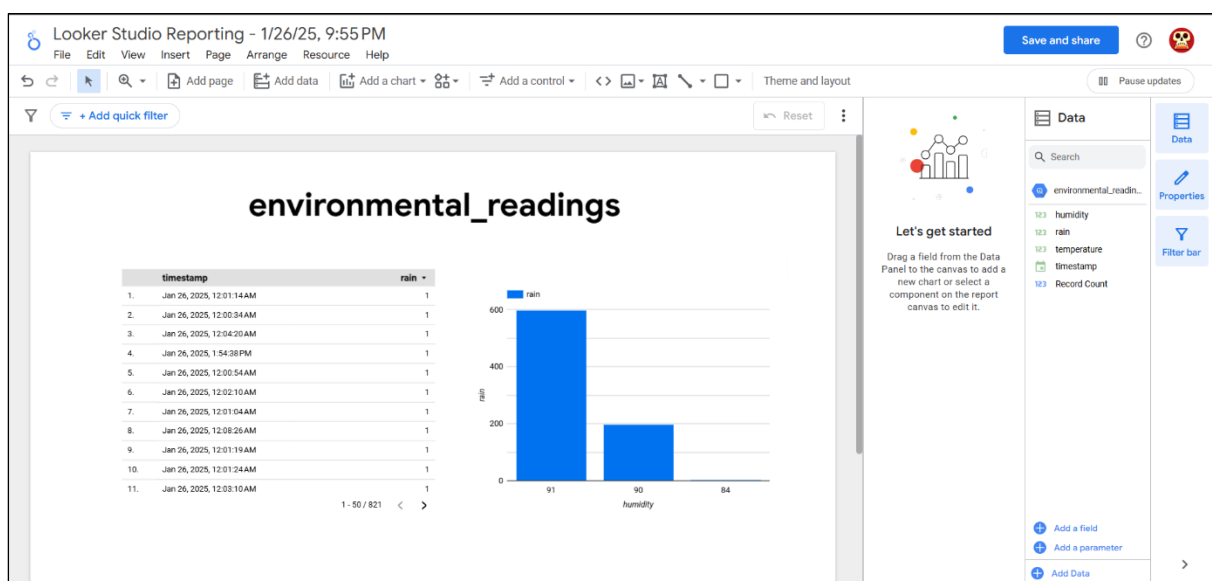
You will see the below output indicating the sensor readings are successfully inserted into BigQuery table.

```
ssh.cloud.google.com/v2/ssh/projects/cpc357-assignment-448819/zones/us-central1-a/instances/instance-20250124-195503?authuser=0&hl=en_US&projectNumber=574179815363&useAdminProxy=true - Google Chrome
ssh.cloud.google.com/v2/ssh/projects/cpc357-assignment-448819/zones/us-central1-a/instances/instance-20250124-195503?authuser=0&hl=en_US&projectNumber=574179815363&useAdminProxy=true
SSH-in-browser
mugu221002@instance-20250124-195503:~/iot-project$ nano mongo-bq.py
mugu221002@instance-20250124-195503:~/iot-project$ nano mongo-bq.py
mugu221002@instance-20250124-195503:~/iot-project$ nano service-account-key.json
mugu221002@instance-20250124-195503:~/iot-project$ ls
mongo-bq.py  service-account-key.json
mugu221002@instance-20250124-195503:~/iot-project$ nano mongo-bq.py
mugu221002@instance-20250124-195503:~/iot-project$ python3 mongo-bq.py
Starting Kqrr listener and BigQuery uploader...
Successfully connected to MQTT broker
Received message: Temperature: 28.50 °C, Humidity: 84.00 %, Rain: 0
Row successfully inserted into BigQuery
Received message: Temperature: 28.50 °C, Humidity: 84.00 %, Rain: 1
Row successfully inserted into BigQuery
Received message: Temperature: 28.50 °C, Humidity: 84.00 %, Rain: 0
Row successfully inserted into BigQuery
Received message: Temperature: 28.50 °C, Humidity: 84.00 %, Rain: 1
Row successfully inserted into BigQuery
Received message: Temperature: 28.50 °C, Humidity: 84.00 %, Rain: 0
Row successfully inserted into BigQuery
```

Now navigate to the BigQuery table and click the **Export > Explore with Looker Studio**. You will be redirect to Looker Studio where you can create charts according to your preference.

The screenshot shows the Google Cloud BigQuery Explorer interface. The left sidebar displays the project hierarchy, including 'cpc357-assignment-448819' and 'sensor_data'. The main area shows the 'environmental_readings' table with a data preview. The 'EXPORT' button is highlighted, and a dropdown menu is open, showing options: 'Explore with Sheets', 'Explore with Looker Studio', 'Export to GCS', and 'Scan with Sensitive Data Protection'.

Row	timestamp	temperature	humidity	rain
1	2025-01-26 00:00:29.656241 U...	28.0	90.0	1
2	2025-01-26 00:00:34.215858 U...	28.0	90.0	1
3	2025-01-26 00:00:39.244039 U...	28.0	90.0	1
4	2025-01-26 00:00:44.263843 U...	28.0	90.0	1
5	2025-01-26 00:00:49.723895 U...	28.0	90.0	1
6	2025-01-26 00:00:54.316004 U...	28.0	90.0	1
7	2025-01-26 00:00:59.253945 U...	28.0	90.0	1
8	2025-01-26 00:01:04.373986 U...	28.0	90.0	1
9	2025-01-26 00:01:09.794087 U...	28.0	90.0	1
10	2025-01-26 00:01:14.406076 U...	28.0	90.0	1
11	2025-01-26 00:01:19.454240 U...	28.0	90.0	1
12	2025-01-26 00:01:24.465875 U...	28.0	90.0	1
13	2025-01-26 00:01:29.864977 U...	28.0	90.0	1
14	2025-01-26 00:01:34.526043 U...	28.0	90.0	1
15	2025-01-26 00:01:39.555855 U...	28.0	90.0	1
16	2025-01-26 00:01:44.573777 U...	28.0	90.0	1
17	2025-01-26 00:01:49.933777 U...	28.0	90.0	1
18	2025-01-26 00:01:54.624914 U...	28.0	90.0	1



5. Security Considerations

5.1 Security Measures Implemented

1. Transport Layer Security (TLS) for MQTT Communication:

All communications between the **microcontroller** and the **MQTT broker** are encrypted using **TLS** over port **8883**, ensuring that data remains **confidential** and protected from **eavesdropping** or **man-in-the-middle attacks** during transmission. This encryption mechanism guarantees that even if data packets are intercepted, they cannot be read or altered without proper decryption keys. Additionally, the use of **TLS certificates** validates both the client and server, establishing a trusted connection and preventing communication with unauthorized or rogue servers.

2. Firewall Rules in Virtual Private Cloud (VPC):

Firewall rules within the **Virtual Private Cloud (VPC)** are configured to regulate **inbound** and **outbound traffic**, ensuring that only **authorized devices** and **IP ranges** can interact with the **MQTT broker** and other system resources. Specific ports are selectively opened to limit exposure to vulnerabilities:\n - **Port 8883** is enabled exclusively for encrypted MQTT traffic.

- **Port 22**, used for SSH access, is further secured by restricting access to specific **IP addresses**. This strict control minimizes the system's attack surface and prevents unauthorized attempts to access cloud components, ensuring **network-level security**.

3. Identity and Access Management (IAM):

The system leverages **Google Cloud IAM** to implement **role-based access control (RBAC)**, ensuring that each service and user has the **minimum necessary permissions** required to perform their functions. This adherence to the principle of **least privilege** minimizes the risk of accidental or malicious misuse of resources. Examples include:

- Assigning the **BigQuery Admin role** exclusively to service accounts responsible for **managing data storage** and **running queries**.
- Restricting **end-user accounts** to **read-only permissions** for dashboards in **Looker Studio**, ensuring they cannot modify or delete data. IAM policies are reviewed

regularly to align with changing system requirements, enhancing security over time.

4. **Authentication and Authorization:**

Secure **API keys** and **OAuth tokens** are used to authenticate connections between the microcontroller, MQTT broker, and cloud services. These tokens ensure that only **verified devices** and **users** can access the system's endpoints. Furthermore, API keys are periodically rotated, and token lifetimes are limited to prevent misuse in case of a compromise. All **API endpoints** are configured to reject unauthenticated requests, providing a robust framework to prevent **unauthorized access**.

5. **Data Encryption in Transit and at Rest:**

Data is encrypted **in transit** using **TLS** for MQTT communication and **HTTPS** for accessing dashboards. This encryption prevents the exposure of sensitive data as it moves between system components. Additionally, all data stored in **BigQuery** and **MongoDB** is encrypted **at rest** using Google Cloud's default **AES-256 encryption**. This ensures that even if storage devices are physically compromised, the data remains secure and unreadable without the decryption keys. Google Cloud's encryption mechanisms are compliant with industry standards, providing **trust and reliability** to users.

5.2 Strategies for Data Protection and Access Control

Regular Key Rotation

To minimize the risk of compromised credentials, the system implements **regular key rotation** for all **encryption keys** and **API keys**. This practice ensures that even if a key is exposed, its usage window is limited. By automating this process, the system reduces the administrative burden and enhances the overall **security posture**. This proactive approach helps maintain **data confidentiality** and limits exposure to potential threats.

Access Control Policies

The system enforces **granular access policies** to ensure that only **authorized users** and devices have access to specific resources. **Role-based access control (RBAC)** is implemented using Google Cloud **Identity and Access Management (IAM)**, assigning roles with **minimum permissions** required for tasks. For example, the **BigQuery Admin role** allows managing data storage, while **end-user accounts** have read-only access to visualizations in **Looker Studio**. These policies are crucial for maintaining **system integrity** and preventing unauthorized access.

Monitoring and Alerts

Real-time monitoring is set up through **Google Cloud Monitoring**, which tracks system activity and identifies unusual patterns. Alerts are configured to notify administrators about **potential breaches**, such as unauthorized access attempts or anomalies in data traffic. By reviewing logs from **Cloud Audit Logs**, administrators can quickly identify and respond to suspicious activities. This continuous monitoring ensures **timely interventions** to mitigate risks and maintain **data security**.

Backup and Disaster Recovery Plan

The system employs a **robust backup strategy**, regularly saving critical data to secure locations. This ensures data can be restored quickly in case of a breach or system failure. Additionally, a **disaster recovery plan** is in place to restore services within **minutes**, ensuring minimal downtime. These measures protect the system against **data loss** and provide operational resilience during **unexpected events**.

Penetration Testing and Vulnerability Scanning

The system undergoes periodic **penetration testing** to identify and fix vulnerabilities that could be exploited by attackers. Automated **vulnerability scans** are also conducted to detect **misconfigurations**, outdated libraries, or exposed endpoints. By addressing these weaknesses proactively, the system ensures **data integrity** and reduces the likelihood of successful cyberattacks.

Token Expiration and Session Management

Authentication tokens are configured with short expiration times, requiring users and devices to regularly **reauthenticate**. This reduces the risk of unauthorized access if tokens are intercepted. Additionally, session timeouts are enforced, automatically logging out inactive users to prevent **session hijacking**. These measures enhance **access control** while maintaining a balance between **usability** and **security**.

6. Appendix

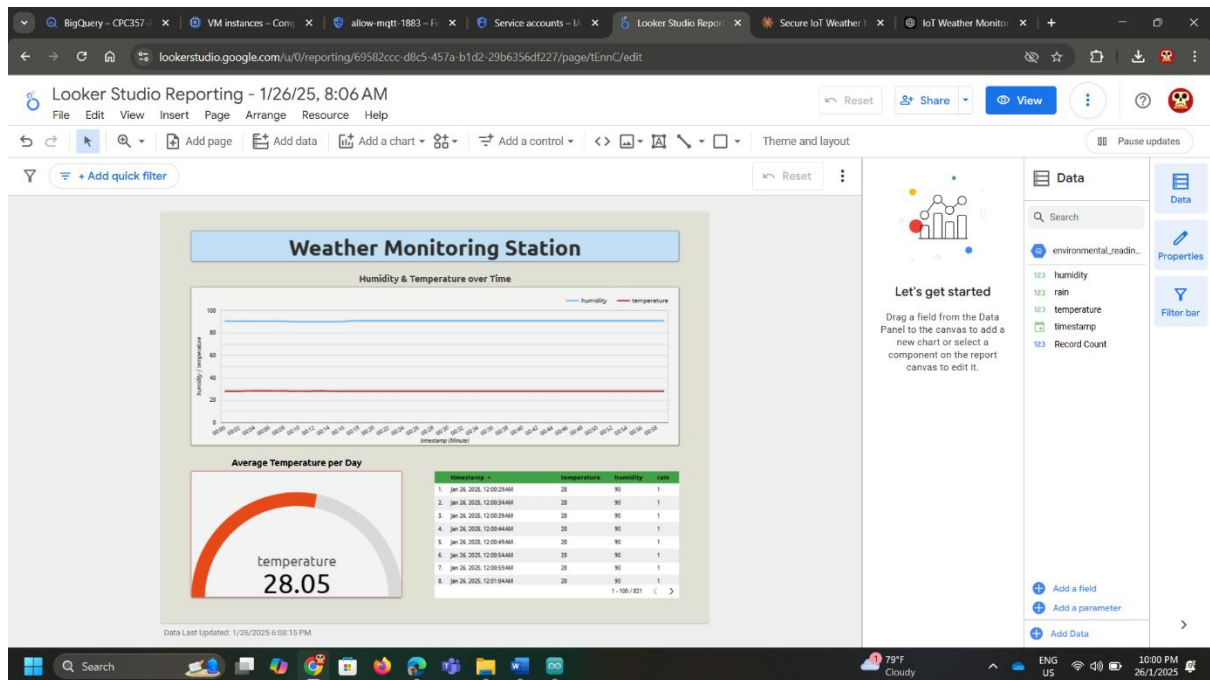


Figure 1A: Dashboard of the Weather Monitoring Station

GitHub Repo: <https://github.com/mugunthantypical/CPC357-Assignment-2.git>